

```
In [ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import plotly.express as px
%matplotlib inline
```

```
In [ ]: pd.set_option('display.max_rows', 10)
pd.options.display.max_columns = 100
pd.set_option("display.precision", 2)
```

Data Loading and cleaning

Initial Analysis of the Datasets

1. Income inequality

```
In [ ]: income = pd.read_csv('../CSV/income_per_person_gdppercapita_ppp_inflation_acc')
```

```
In [ ]: print("number of rows: ", income.shape[0])
print("number of columns: {}".format(income.shape[1]))
print("number of duplicates: {}".format(income.duplicated().sum()))
print("datatypes:\n")
print(income.dtypes)
income.head(3)
```

```
number of rows: 193
number of columns: 242
number of duplicates: 0
datatypes:
```

```
country    object
1800      int64
1801      int64
1802      int64
1803      int64
...
2036      int64
2037      int64
2038      int64
2039      int64
2040      int64
Length: 242, dtype: object
```

```
Out[ ]:
```

	country	1800	1801	1802	1803	1804	1805	1806	1807	1808	1809	1810	1811
0	Afghanistan	603	603	603	603	603	603	603	603	603	603	604	604
1	Albania	667	667	667	667	667	668	668	668	668	668	668	668
2	Algeria	715	716	717	718	719	720	721	722	723	724	725	726

3 rows x 242 columns

The results below show that there no nulls.

```
In [ ]: income_null = income.isnull().sum()/income.shape[0]
income_null.to_frame().transpose()
```

```
Out[ ]:    country 1800 1801 1802 1803 1804 1805 1806 1807 1808 1809 1810 1811 1812
0      0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
```

1 rows × 242 columns

2. Tax Revenue as a Percent of GDP

```
In [ ]: tax = pd.read_csv('../CSV/tax_revenue_percent_of_gdp.csv')
```

```
In [ ]: print("number of rows: ", tax.shape[0])
print("number of columns: {}".format(tax.shape[1]))
print("number of duplicates: {}".format(tax.duplicated().sum()))
print("datatypes:\n")
print(tax.dtypes)
print("\nSample:")
tax.head(3)
```

```
number of rows: 161
number of columns: 47
number of duplicates: 0
datatypes:
```

```
country      object
1972         float64
1973         float64
1974         float64
1975         float64
...
2013         float64
2014         float64
2015         float64
2016         float64
2017         float64
Length: 47, dtype: object
```

Sample:

```
Out[ ]:    country 1972 1973 1974 1975 1976 1977 1978 1979 1980 1981 1982 1983
0  Afghanistan  NaN  NaN  NaN  NaN  NaN  NaN  NaN  NaN  NaN  NaN  NaN
1    Albania    NaN  NaN  NaN  NaN  NaN  NaN  NaN  NaN  NaN  NaN  NaN
2    Algeria    NaN  NaN  NaN  NaN  NaN  NaN  NaN  NaN  NaN  NaN  NaN
```

An initial analysis revealed that about half the years with > 0.50 missing values. See below:

```
In [ ]: tax_null = tax.isnull().sum()/tax.shape[0]
tax_null.to_frame().transpose()
```

```
Out [ ]:      country 1972 1973 1974 1975 1976 1977 1978 1979 1980 1981 1982 1983 19
0         0.0  0.76  0.72  0.72  0.71  0.71  0.7  0.7  0.7  0.68  0.68  0.7  0.69  0
```

3. Gini dataset

```
In [ ]: gini = pd.read_csv('../CSV/gini.csv')
```

```
In [ ]: print("number of rows: ", gini.shape[0])
print("number of columns: {}".format(gini.shape[1]))
print("number of duplicates: {}".format(gini.duplicated().sum()))
print("datatypes:\n")
print(gini.dtypes)
print("\nSample:")
gini.head(3)
```

```
number of rows: 195
number of columns: 242
number of duplicates: 0
datatypes:
```

```
country      object
1800         float64
1801         float64
1802         float64
1803         float64
...
2036         float64
2037         float64
2038         float64
2039         float64
2040         float64
Length: 242, dtype: object
```

Sample:

```
Out [ ]:      country 1800 1801 1802 1803 1804 1805 1806 1807 1808 1809 1810 1811
0  Afghanistan  30.5  30.5  30.5  30.5  30.5  30.5  30.5  30.5  30.5  30.5  30.5  30.5
1    Albania    38.9  38.9  38.9  38.9  38.9  38.9  38.9  38.9  38.9  38.9  38.9  38.9
2    Algeria    56.2  56.2  56.2  56.2  56.2  56.2  56.2  56.2  56.2  56.2  56.2  56.2
```

3 rows x 242 columns

The results below show that there were no nulls in the dataset.

```
In [ ]: gini_null = gini.isnull().sum()/gini.shape[0]
gini_null.to_frame().transpose()
```

```
Out [ ]:      country 1800 1801 1802 1803 1804 1805 1806 1807 1808 1809 1810 1811 1812
```

	country	1800	1801	1802	1803	1804	1805	1806	1807	1808	1809	1810	1811	1812
0		0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

1 rows x 242 columns

4. Investment Percent of GDP

Below are results of an initial analysis:

```
In [ ]: invest = pd.read_csv('../CSV/investments_percent_of_gdp.csv')
```

```
In [ ]: print("number of rows: ", invest.shape[0])
print("number of columns: {}".format(invest.shape[1]))
print("number of duplicates: {}".format(invest.duplicated().sum()))
print("datatypes:\n")
print(invest.dtypes)
print("\nSample:")
invest.head(3)
```

number of rows: 177
 number of columns: 59
 number of duplicates: 0
 datatypes:

```
country      object
1960         float64
1961         float64
1962         float64
1963         float64
...
2013         float64
2014         float64
2015         float64
2016         float64
2017         float64
Length: 59, dtype: object
```

Sample:

```
Out [ ]:      country 1960 1961 1962 1963 1964 1965 1966 1967 1968 1969 1970 1971 1972
```

	country	1960	1961	1962	1963	1964	1965	1966	1967	1968	1969	1970	1971	1972
0	Afghanistan	16.1	16.6	19.1	14.2	13.9	11.3	8.41	5.18	6.47	6.47	5.46	5.46	5.46
1	Albania	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2	Algeria	42.2	47.2	35.4	28.9	21.8	22.6	17.30	23.40	27.90	32.40	36.60	35.40	35.40

```
In [ ]: invest_null = invest.isnull().sum()/invest.shape[0]
invest_null.to_frame().transpose()
```

```
Out [ ]:      country 1960 1961 1962 1963 1964 1965 1966 1967 1968 1969 1970 1971 1972
```

	country	1960	1961	1962	1963	1964	1965	1966	1967	1968	1969	1970	1971	1972
0		0.0	0.63	0.63	0.63	0.62	0.54	0.53	0.51	0.51	0.51	0.44	0.44	0.44

5. EIU Democracy Index

Below are the results of the initial analysis:

```
In [ ]: demo = pd.read_csv('../CSV/demox_eiu.csv')
```

```
In [ ]: print("number of rows: ", demo.shape[0])
print("number of columns: {}".format(demo.shape[1]))
print("number of duplicates: {}".format(demo.duplicated().sum()))
print("datatypes:\n")
print(demo.dtypes)
print("\nSample:")
demo.head(3)
```

```
number of rows: 164
number of columns: 14
number of duplicates: 0
datatypes:
```

```
country      object
2006         float64
2007         float64
2008         float64
2009         float64
...
2014         float64
2015         float64
2016         float64
2017         float64
2018         float64
Length: 14, dtype: object
```

Sample:

```
Out [ ]: 
```

	country	2006	2007	2008	2009	2010	2011	2012	2013	2014	2015	2016	2017
0	Afghanistan	30.6	30.4	30.2	27.5	24.8	24.8	24.8	24.8	27.7	27.7	25.5	25.5
1	Albania	59.1	59.1	59.1	58.9	58.6	58.1	56.7	56.7	56.7	59.1	59.1	59.1
2	Algeria	31.7	32.5	33.2	33.8	34.4	34.4	38.3	38.3	38.3	39.5	35.6	35.6

The results below show that there were no nulls in the dataset.

```
In [ ]: demo_null = demo.isnull().sum()/demo.shape[0]
demo_null.to_frame().transpose()
```

```
Out [ ]: 
```

	country	2006	2007	2008	2009	2010	2011	2012	2013	2014	2015	2016	2017
0		0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

Because of the results below, the years with the lowest null value percentage were in the 10 year range from 2006-2016. As a result, this yearly range was selected for the rest of the datasets.

Slicing and Reorganizing the Datasets

Income Per Person (GDP per Capita)

```
In [ ]: income_last_10 = income.iloc[:, np.r_[1, 207:218]]
income_last_10
```

```
Out[ ]:
```

	country	2006	2007	2008	2009	2010	2011	2012	2013	2014	2015
0	Afghanistan	1120	1250	1270	1500	1670	1630	1770	1810	1800	1770
1	Albania	7910	8450	9160	9530	9930	10200	10400	10500	10700	11000
2	Algeria	12400	12600	12700	12700	12900	13000	13200	13300	13500	13800
3	Andorra	42700	43400	41400	41700	39000	42000	41900	43700	44900	46600
4	Angola	5500	6040	6470	6290	6360	6350	6650	6730	6810	6650
...
188	Venezuela	16400	17600	18200	17400	16900	17300	18000	18000	17100	15600
189	Vietnam	3630	3850	4030	4210	4430	4660	4860	5070	5310	5610
190	Yemen	4270	4290	4320	4360	4570	3880	3860	3940	3830	3110
191	Zambia	2650	2800	2930	3120	3340	3420	3570	3630	3690	3680
192	Zimbabwe	1890	1810	1480	1630	1930	2170	2490	2490	2510	2510

193 rows x 12 columns

```
In [ ]: income_last_10 = income_last_10.melt(id_vars=['country'], var_name='year', \
income_last_10.sort_values(['country', 'year'], inplace=True)
```

```
In [ ]: income_last_10.head(3)
```

```
Out[ ]:
```

	country	year	income_per_person
0	Afghanistan	2006	1120
193	Afghanistan	2007	1250
386	Afghanistan	2008	1270

Investment Percent of GDP

```
In [ ]: invest_last_10 = invest.iloc[:, np.r_[1, 47:58]]
invest_last_10 = invest_last_10.melt(id_vars=['country'], var_name='year', \
invest_last_10.sort_values(['country', 'year'], inplace=True)
invest_last_10.head(3)
```

```
Out[ ]:
```

	country	year	invest_%_gdp
0	Afghanistan	2006	23.4
177	Afghanistan	2007	19.9
354	Afghanistan	2008	18.9

Tax Revenue Percent of GDP

```
In [ ]: tax_last_10 = tax.iloc[:, np.r_[:1, 35:46]]
tax_last_10 = tax_last_10.melt(id_vars=['country'], var_name='year', value_
tax_last_10.sort_values(['country', 'year'], inplace=True)
tax_last_10.head(3)
```

```
Out[ ]:
```

	country	year	tax_%_gdp
0	Afghanistan	2006	6.88
161	Afghanistan	2007	5.23
322	Afghanistan	2008	6.04

Gini Index

```
In [ ]: gini_last_10 = gini.iloc[:, np.r_[:1, 207:218]]
gini_last_10 = gini_last_10.melt(id_vars=['country'], var_name='year', value_
gini_last_10.sort_values(by=['country', 'year'], inplace=True)
gini_last_10.head(3)
```

```
Out[ ]:
```

	country	year	gini_index
0	Afghanistan	2006	36.8
195	Afghanistan	2007	36.8
390	Afghanistan	2008	36.8

EIU Democracy Index

```
In [ ]: demo_last_10 = demo.iloc[:, :-2]
demo_last_10 = demo_last_10.melt(id_vars=['country'], var_name='year', value_
demo_last_10.sort_values(['country', 'year'], inplace=True)
demo_last_10.head(3)
```

```
Out[ ]:
```

	country	year	demox_eiu
0	Afghanistan	2006	30.6
164	Afghanistan	2007	30.4
328	Afghanistan	2008	30.2

Merging the Datasets

```
In [ ]: combined = demo_last_10.merge(income_last_10, left_on=['country', 'year'],
combined = combined.merge(invest_last_10, left_on=['country', 'year'], right_
combined = combined.merge(tax_last_10, left_on=['country', 'year'], right_on
combined = combined.merge(gini_last_10, left_on=['country', 'year'], right_
combined
```

Out []:

	country	year	demox_eiu	income_per_person	invest_%_gdp	tax_%_gdp	gini_ir
0	Afghanistan	2006	30.6	1120	23.4	6.88	
1	Afghanistan	2007	30.4	1250	19.9	5.23	
2	Afghanistan	2008	30.2	1270	18.9	6.04	
3	Afghanistan	2009	27.5	1500	17.9	8.44	
4	Afghanistan	2010	24.8	1670	17.9	9.12	
...
1524	Zimbabwe	2012	26.7	2490	11.8	21.40	
1525	Zimbabwe	2013	26.7	2490	11.4	NaN	
1526	Zimbabwe	2014	27.8	2510	11.8	NaN	
1527	Zimbabwe	2015	30.5	2510	12.3	NaN	
1528	Zimbabwe	2016	30.5	2490	12.2	NaN	

1529 rows × 7 columns

In []: `cont = pd.read_csv('../CSV/continent_country.csv')`
`cont`

Out []:

	continent	country
0	Africa	Congo, Dem. Rep.
1	Africa	Congo, Rep.
2	Africa	Algeria
3	Africa	Angola
4	Africa	Benin
...
165	Europe	Ukraine
166	Europe	United Kingdom
167	Oceania	Australia
168	Oceania	New Zealand
169	continent	country

170 rows × 2 columns

Matching Country with Continent

In this step, we match each country with its continent. This will enable analysis at the continent level for broader trend detection.

In []: `combined_final = cont.merge(combined, left_on=['country'], right_on=['country'])`
`combined_final`

Out[]:

	continent	country	year	demox_eiu	income_per_person	invest_%_gdp	tax_%_gdp
0	Africa	Congo, Dem. Rep.	2006	27.6	605	14.6	6.83
1	Africa	Congo, Dem. Rep.	2007	25.2	623	13.7	6.99
2	Africa	Congo, Dem. Rep.	2008	22.8	640	10.9	8.97
3	Africa	Congo, Dem. Rep.	2009	22.1	637	14.6	7.89
4	Africa	Congo, Dem. Rep.	2010	21.5	660	28.8	8.35
...
1524	Oceania	New Zealand	2012	92.6	33300	20.9	26.80
1525	Oceania	New Zealand	2013	92.6	33900	22.0	26.60
1526	Oceania	New Zealand	2014	92.6	34600	22.9	26.80
1527	Oceania	New Zealand	2015	92.6	35200	23.4	27.40
1528	Oceania	New Zealand	2016	92.6	35700	24.4	27.20

1529 rows x 8 columns

Data Cleaning

Below are the steps taken to ensure quality of the dataset:

Missing Values

Below are is a summary of missing values (nulls) in the dataset:

```
In [ ]: combined_final.isna().sum()
```

```
Out[ ]: continent      0
country      0
year         0
demox_eiu    0
income_per_person  0
invest_%_gdp  31
tax_%_gdp    285
gini_index   0
dtype: int64
```

One option for handling the missing 'tax_%_gdp' values would be to replace them with the country's mean. However, some of the countries have all nulls and some have mostly

nulls for this column.

A second option is to drop the rows with nulls. In the interest of simplicity, we will use this option.

```
In [ ]: combined_final.dropna(inplace=True)
combined_final.isna().sum()
```

```
Out[ ]: continent      0
country      0
year         0
demox_eiu    0
income_per_person  0
invest_%_gdp  0
tax_%_gdp    0
gini_index   0
dtype: int64
```

Duplicates

There are no duplicates in the dataset:

```
In [ ]: combined_final.duplicated().sum()
```

```
Out[ ]: 0
```

Descriptive Statistics

Below are descriptive statistics of the dataset. A review of the values indicates that the min, max and mean values appear to be reasonable.

```
In [ ]: combined_final.describe()
```

```
Out[ ]:
```

	demox_eiu	income_per_person	invest_%_gdp	tax_%_gdp	gini_index
count	1234.00	1234.00	1234.00	1234.00	1234.00
mean	62.08	20520.39	24.39	16.77	38.04
std	20.57	19421.17	7.96	7.63	8.39
min	16.60	605.00	0.00	0.04	24.40
25%	46.25	5935.00	19.50	12.33	32.10
50%	64.90	13700.00	23.20	15.90	36.00
75%	78.80	31175.00	27.70	21.30	42.58
max	99.30	120000.00	67.90	62.90	63.90

Save the Cleaned Dataset

```
In [ ]: combined_final.to_csv('combined_final_last_10_years.csv', index=False)
```

Exploratory Data Analysis

Research Question 1 - Is Income Inequality Getting Worse or Better in the Last 10 Years?

Better means the Gini Index is going down.

Global Gini Mean By Year

```
In [ ]: columns = ['year', 'gini_index']
gini = combined_final[columns]
gini
```

```
Out[ ]:
```

	year	gini_index
0	2006	42.2
1	2007	42.1
2	2008	42.1
3	2009	42.1
4	2010	42.1
...
1524	2012	33.5
1525	2013	34.0
1526	2014	34.0
1527	2015	34.5
1528	2016	34.8

1234 rows × 2 columns

```
In [ ]: gini_annual_average = gini.groupby('year')['gini_index'].mean()
gini_annual_average
```


```
Out[ ]:
```

year	
2006	38.67
2007	38.65
2008	38.53
2009	38.51
2010	38.29
...	...
2012	38.05
2013	37.48
2014	37.32
2015	37.43
2016	37.32

Name: gini_index, Length: 11, dtype: float64

As the plot below shows, the mean global gini index has been going down over the last 10 years, meaning global income inequality is improving.

```
In [ ]: plt.plot(gini_annual_average.index, gini_annual_average)
plt.title('Mean Global Gini Index by Year')
plt.xlabel('Year')
plt.ylabel('Mean Global Gini Index');
```



Mean Global Gini Index by Continent:

```
In [ ]: columns = ['year', 'continent', 'gini_index']
        gini = combined_final[columns]
        gini
```

```
Out[ ]:
```

	year	continent	gini_index
0	2006	Africa	42.2
1	2007	Africa	42.1
2	2008	Africa	42.1
3	2009	Africa	42.1
4	2010	Africa	42.1
...
1524	2012	Oceania	33.5
1525	2013	Oceania	34.0
1526	2014	Oceania	34.0
1527	2015	Oceania	34.5
1528	2016	Oceania	34.8


1234 rows × 3 columns

```
In [ ]: gini_cont_average = gini.groupby(['year', 'continent'])['gini_index'].mean()
        gini_cont_average
```

```
Out[ ]: year continent
2006 Africa 43.01
Americas 49.15
Asia 38.13
Europe 31.73
Oceania 33.75
...
2016 Africa 46.67
Americas 44.81
Asia 35.25
Europe 31.50
Oceania 33.55
Name: gini_index, Length: 55, dtype: float64
```

The chart below reveals that, on a continent basis, all were either declining or mostly flat, except for Africa.

```
In [ ]: gini_cont_average.unstack(level=1).plot(kind='line', subplots=False, \
                                                title='Mean Global Gini Index by Con\
                                                set_ylabel("Gini Index");
```



```
In [ ]: gini_cont_average.unstack(level=1)
```

Out []: **continent** **Africa** **Americas** **Asia** **Europe** **Oceania**

year					
2006	43.01	49.15	38.13	31.73	33.75
2007	43.42	48.76	37.92	31.62	33.80
2008	43.78	48.36	37.31	31.67	33.80
2009	43.80	48.00	37.22	31.58	33.65
2010	43.54	47.58	36.40	31.58	33.40
...
2012	44.08	46.96	36.33	31.61	32.95
2013	44.07	45.78	36.26	31.59	32.95
2014	44.82	45.48	35.65	31.61	33.10
2015	45.54	45.29	35.54	31.56	33.40
2016	46.67	44.81	35.25	31.50	33.55

11 rows × 5 columns

```
In [ ]: columns = ['year', 'continent', 'country', 'gini_index']
gini = combined_final[columns]
gini
```

Out []: **year** **continent** **country** **gini_index**

0	2006	Africa	Congo, Dem. Rep.	42.2
1	2007	Africa	Congo, Dem. Rep.	42.1
2	2008	Africa	Congo, Dem. Rep.	42.1
3	2009	Africa	Congo, Dem. Rep.	42.1
4	2010	Africa	Congo, Dem. Rep.	42.1
...
1524	2012	Oceania	New Zealand	33.5
1525	2013	Oceania	New Zealand	34.0
1526	2014	Oceania	New Zealand	34.0
1527	2015	Oceania	New Zealand	34.5
1528	2016	Oceania	New Zealand	34.8

1234 rows × 4 columns

Research Question 2 - What Top 10 Countries Have the Lowest and Highest Income Inequality?

Lowest

Overall, most of the countries with the lowest income inequality are in Europe.

```
In [ ]: gini.groupby(['country', 'continent'])['gini_index'].mean().to_frame().sort_
```

Out []:

gini_index		
country	continent	
Slovenia	Europe	25.06
Ukraine	Europe	25.64
Czech Republic	Europe	26.25
Norway	Europe	26.75
Slovak Republic	Europe	26.79
Denmark	Europe	27.16
Kazakhstan	Asia	27.44
Finland	Europe	27.45
Belarus	Europe	27.49
Kyrgyz Republic	Asia	27.65

Highest

Overall, most of the countries with the highest income inequality are in Africa and in Americas.

```
In [ ]: gini.groupby(['country', 'continent'])['gini_index'].mean().to_frame().sort_
```

Out []:

gini_index		
country	continent	
South Africa	Africa	63.35
Botswana	Africa	61.09
Namibia	Africa	60.75
Suriname	Americas	60.51
Zambia	Africa	56.44
Central African Republic	Africa	55.70
Bolivia	Americas	54.55
Honduras	Americas	53.94
Lesotho	Africa	53.93
Colombia	Americas	53.16

Research Question 3 - Is a higher tax revenue as a % of GDP associated with less income inequality?

The hypothesis is that countries with higher tax revenue as % of GDP are associated with lower income inequality. The assumption for this is that higher tax revenues are distributed back to lower economic strata in the form of social benefits. Let's see what the data shows.

```
In [ ]: columns = ['continent', 'country', 'year', 'tax_%_gdp', 'gini_index']
tax = combined_final[columns]
```

tax

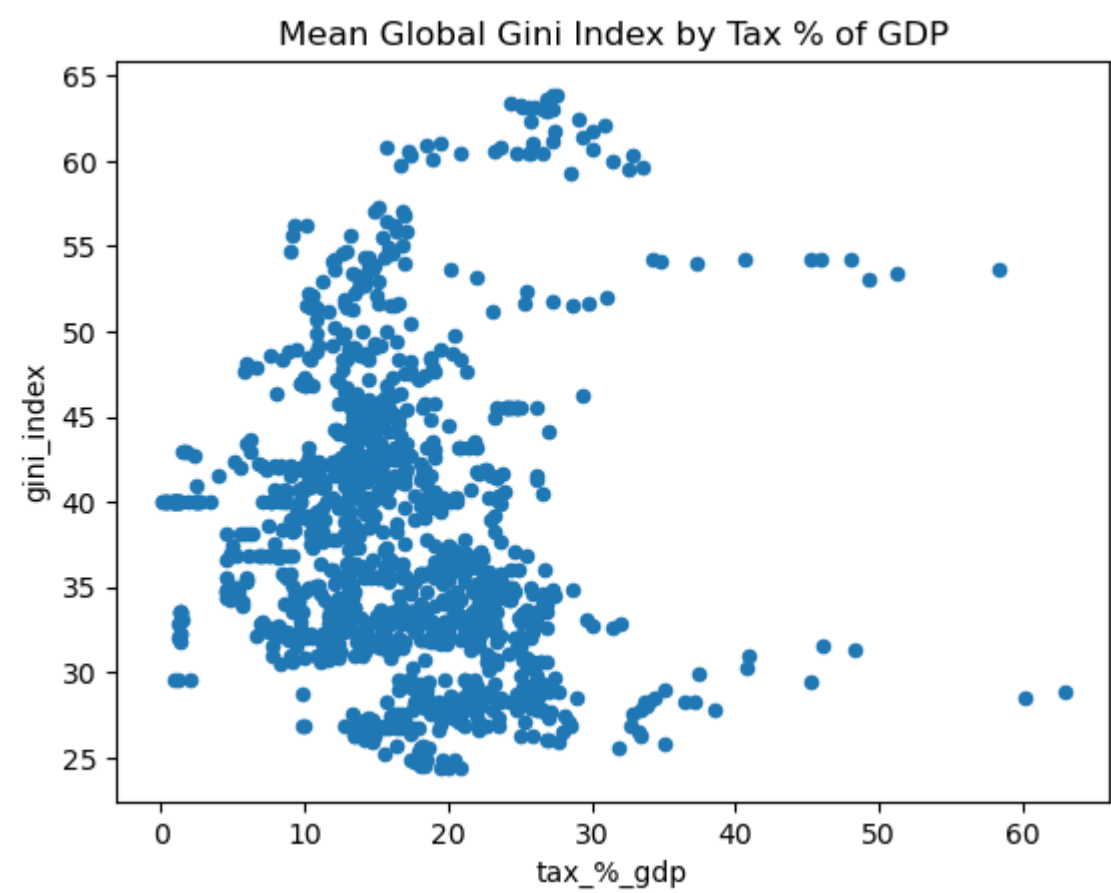
Out[]:

	continent	country	year	tax_%_gdp	gini_index
0	Africa	Congo, Dem. Rep.	2006	6.83	42.2
1	Africa	Congo, Dem. Rep.	2007	6.99	42.1
2	Africa	Congo, Dem. Rep.	2008	8.97	42.1
3	Africa	Congo, Dem. Rep.	2009	7.89	42.1
4	Africa	Congo, Dem. Rep.	2010	8.35	42.1
...
1524	Oceania	New Zealand	2012	26.80	33.5
1525	Oceania	New Zealand	2013	26.60	34.0
1526	Oceania	New Zealand	2014	26.80	34.0
1527	Oceania	New Zealand	2015	27.40	34.5
1528	Oceania	New Zealand	2016	27.20	34.8

1234 rows x 5 columns

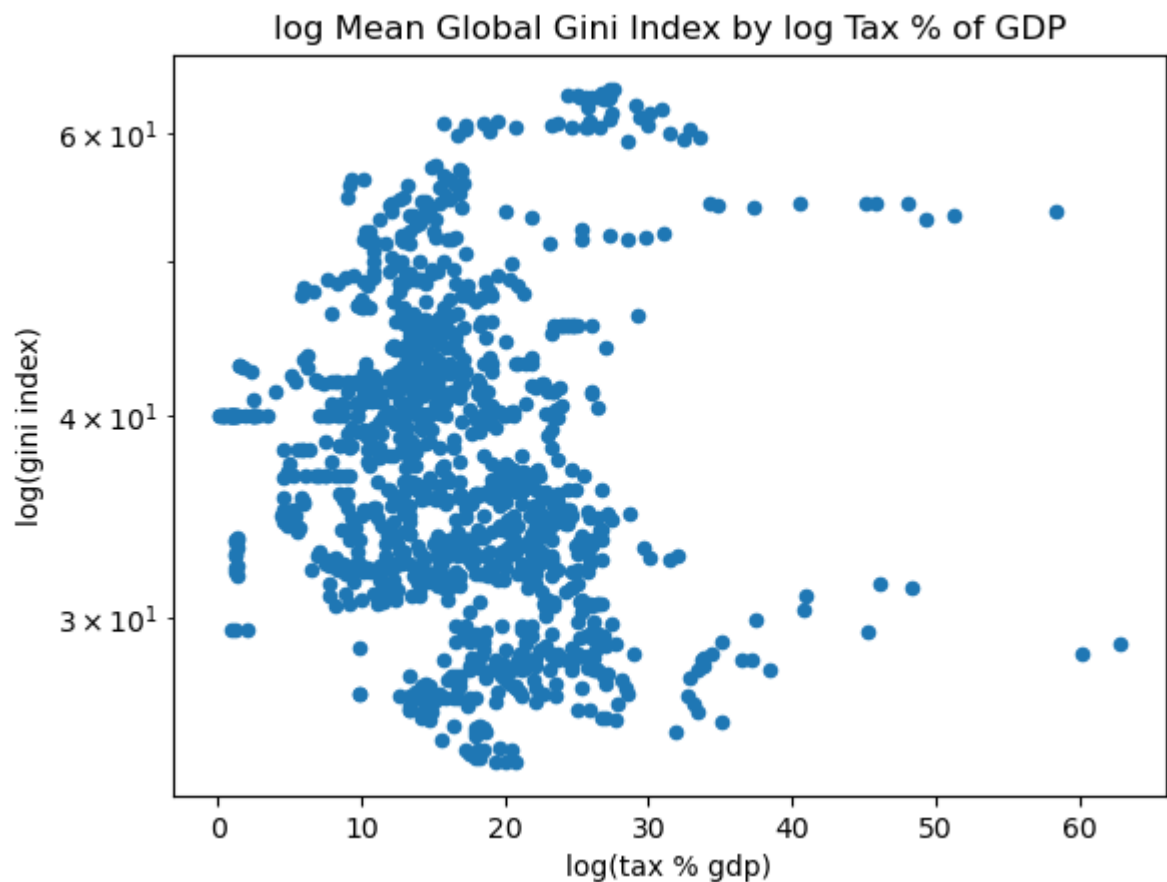
It is difficult to see a trend in the scatter plot below:

```
In [ ]: tax.plot(x='tax_%_gdp', y='gini_index', kind='scatter', title='Mean Global Gini Index by Tax % of GDP')
```



Looking at the log of both values reveals that the correlation between the two variables is essentially flat - there is no compelling evidence that higher tax percent of GDP leads to less income inequality.


```
In [ ]: tax_plot = tax.plot(x='tax_%_gdp', y='gini_index', kind='scatter', logy=True,
                           title='log Mean Global Gini Index by log Tax % of GDP')
tax_plot.set_xlabel('log(tax % gdp)')
tax_plot.set_ylabel('log(gini index)');
```



The Pearson correlation is slightly negative at -0.08:

```
In [ ]: tax_log = np.log(tax['tax_%_gdp']).to_frame()
tax_log['log_gini_index'] = np.log(tax['gini_index'])
tax_log.corr()

tax_corr = tax[['tax_%_gdp', 'gini_index']]
tax_corr.corr()
```

```
Out [ ]:
```

	tax_%_gdp	gini_index
tax_%_gdp	1.00	-0.05
gini_index	-0.05	1.00

Research Question 4 - Is Higher Income Per Person - GDP Per Capita associated with less income inequality?

The hypothesis is that a higher income per person indicates that more of the country's GDP is being distributed equality among its population.

```
In [ ]: columns = ['continent', 'country', 'year', 'income_per_person', 'gini_index']
income = combined_final[columns]
income
```

Out []:


	continent	country	year	income_per_person	gini_index
0	Africa	Congo, Dem. Rep.	2006	605	42.2
1	Africa	Congo, Dem. Rep.	2007	623	42.1
2	Africa	Congo, Dem. Rep.	2008	640	42.1
3	Africa	Congo, Dem. Rep.	2009	637	42.1
4	Africa	Congo, Dem. Rep.	2010	660	42.1
...
1524	Oceania	New Zealand	2012	33300	33.5
1525	Oceania	New Zealand	2013	33900	34.0
1526	Oceania	New Zealand	2014	34600	34.0
1527	Oceania	New Zealand	2015	35200	34.5
1528	Oceania	New Zealand	2016	35700	34.8

1234 rows x 5 columns


```
In [ ]: income.plot(x='income_per_person', y='gini_index', kind='scatter', title='Me
```



```
In [ ]: income_plot = income.plot(x='income_per_person', y='gini_index', kind='scatter',
                                   title='log Mean Gini Index by Income Per Person')
income_plot.set_xlabel('log(income_per_person)')
income_plot.set_ylabel('log(gini index)');
```



```
In [ ]: income_plot = income.plot(x='income_per_person', y='gini_index', kind='scatter',
                                   title='log Mean Gini Index by Income Per Person')
income_plot.set_xlabel('income_per_person')
income_plot.set_ylabel('log(gini index)');
```



In this case, the Person correlation coefficient is -0.34 indicating that there is weak correlation between $\log(\text{income_per_person})$ and the $\log(\text{gini_index})$:

```
In [ ]: income_log = np.log(income['income_per_person']).to_frame()
income_log['log_gini_index'] = np.log(tax['gini_index'])
income_log.columns = ['log_income_per_person', 'log_gini_index']
# income_log.corr()

income_corr = income[['income_per_person', 'gini_index']]
income_corr.corr()
```

```
Out [ ]:
```

	income_per_person	gini_index
income_per_person	1.00	-0.33
gini_index	-0.33	1.00

Research Question 5 - Is Higher Investment as % GDP associated with less income inequality?

The hypothesis is that a higher investment as a percent of GDP indicates that more of the country's GDP is being invested in capital improvements which distributes income benefits across a wide segment of the population leading to more equality among its population.

```
In [ ]: columns = ['continent', 'country', 'year', 'invest_%_gdp', 'gini_index']
invest = combined_final[columns]
invest
```

```
Out [ ]:
```

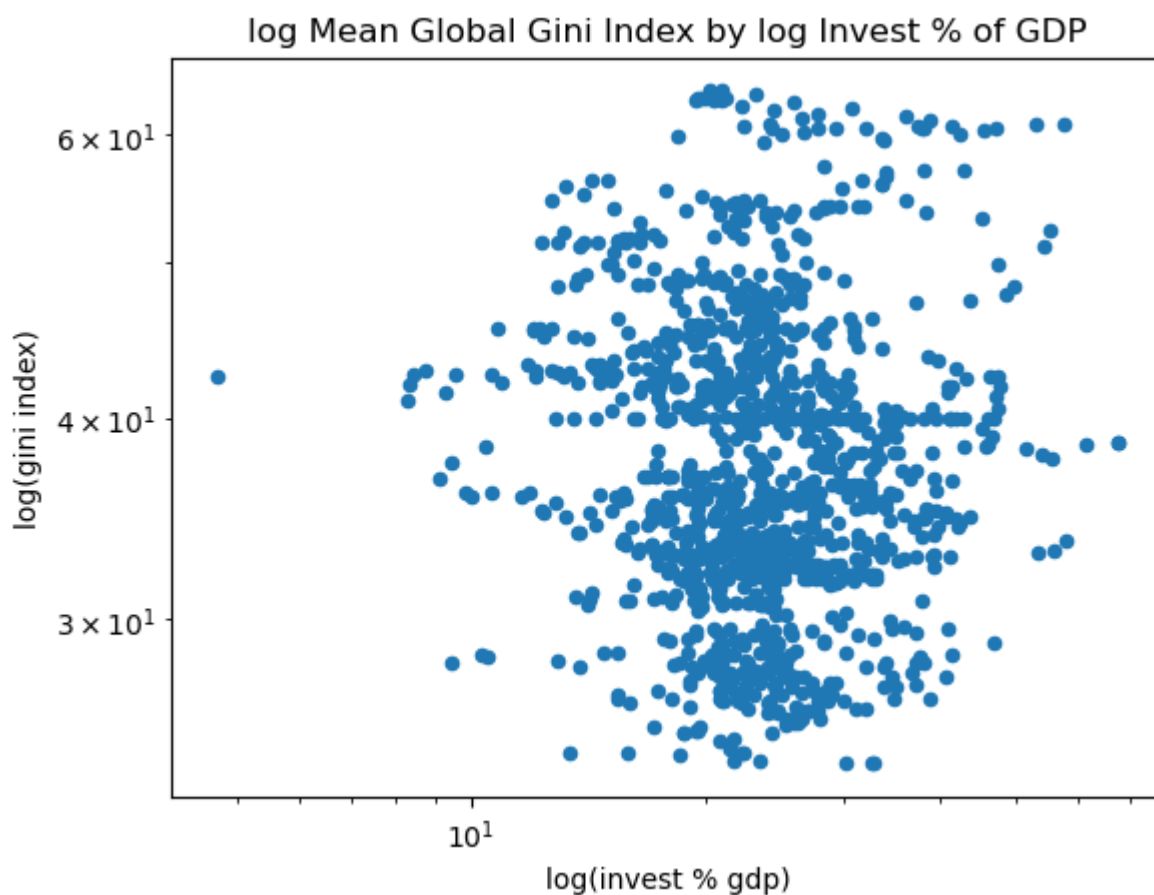
	continent	country	year	invest_%_gdp	gini_index
0	Africa	Congo, Dem. Rep.	2006	14.6	42.2
1	Africa	Congo, Dem. Rep.	2007	13.7	42.1
2	Africa	Congo, Dem. Rep.	2008	10.9	42.1
3	Africa	Congo, Dem. Rep.	2009	14.6	42.1
4	Africa	Congo, Dem. Rep.	2010	28.8	42.1
...
1524	Oceania	New Zealand	2012	20.9	33.5
1525	Oceania	New Zealand	2013	22.0	34.0
1526	Oceania	New Zealand	2014	22.9	34.0
1527	Oceania	New Zealand	2015	23.4	34.5
1528	Oceania	New Zealand	2016	24.4	34.8

1234 rows × 5 columns

```
In [ ]: invest = invest[invest['invest_%_gdp'] > 0]
invest.plot(x='invest_%_gdp', y='gini_index', kind='scatter', title='Mean G:
```



```
In [ ]: invest_plot = invest.plot(x='invest_%_gdp', y='gini_index', kind='scatter',  
                                title='log Mean Global Gini Index by log Invest % of GDP',  
                                invest_plot.set_xlabel('log(invest % gdp)'),  
                                invest_plot.set_ylabel('log(gini index)');
```



The Pearson corr coefficient of -0.03 indicates no correlation between these two variables.

```
In [ ]: invest_log = np.log(invest['invest_%_gdp']).to_frame()
invest_log['log_gini_index'] = np.log(tax['gini_index'])
invest_log.corr()

invest_corr = invest[['invest_%_gdp', 'gini_index']]
invest_corr.corr()
```

```
Out [ ]:
```

	invest_%_gdp	gini_index
invest_%_gdp	1.00	0.02
gini_index	0.02	1.00

Research Question 6 - Is Higher EIU Democracy Index associated with less income inequality?

The hypothesis is that countries with higher EIU Democracy Index address the needs of a broader segment of the population leading to less income inequality.


```
In [ ]: columns = ['continent', 'country', 'year', 'demox_eiu', 'gini_index']
demo = combined_final[columns]
demo
```

```
Out [ ]:
```


	continent	country	year	demox_eiu	gini_index
0	Africa	Congo, Dem. Rep.	2006	27.6	42.2
1	Africa	Congo, Dem. Rep.	2007	25.2	42.1
2	Africa	Congo, Dem. Rep.	2008	22.8	42.1
3	Africa	Congo, Dem. Rep.	2009	22.1	42.1
4	Africa	Congo, Dem. Rep.	2010	21.5	42.1
...
1524	Oceania	New Zealand	2012	92.6	33.5
1525	Oceania	New Zealand	2013	92.6	34.0
1526	Oceania	New Zealand	2014	92.6	34.0
1527	Oceania	New Zealand	2015	92.6	34.5
1528	Oceania	New Zealand	2016	92.6	34.8

1234 rows x 5 columns

```
In [ ]: demo.plot(x='demox_eiu', y='gini_index', kind='scatter', title='Mean Gini I
```



```
In [ ]: demo_plot = demo.plot(x='demox_eiu', y='gini_index', kind='scatter', loglog=
        title='log Mean Global Gini Index by log EIU Democracy Index')
demo_plot.set_xlabel('log(demox_eiu)')
demo_plot.set_ylabel('log(gini_index)');
```



In this case, the Person correlation coefficient is -0.2 indicating that there is weak correlation between $\log(\text{demox_eiu})$ and the $\log(\text{gini_index})$:

```
In [ ]: demo_log = np.log(demo['demox_eiu']).to_frame()
demo_log['log_gini_index'] = np.log(tax['gini_index'])
demo_log.corr()
```

```
Out [ ]:
```

	demox_eiu	log_gini_index
demox_eiu	1.0	-0.2
log_gini_index	-0.2	1.0

Conclusions

The following are the conclusions from this analysis:

Research Question 1 - Is Income Inequality Getting Worse or Better in the Last 10 Years?

Answer:

Yes, it is getting better, improving from 38.7 to 37.3

On a continent basis, all were either declining or mostly flat, except for Africa.

Research Question 2 - What Top 10 Countries Have the Lowest and Highest Income Inequality?

Answer:

Lowest: Slovenia, Ukraine, Czech Republic, Norway, Slovak Republic, Denmark, Kazakhstan, Finland, Belarus, Kyrgyz Republic

Highest: Colombia, Lesotho, Honduras, Bolivia, Central African Republic, Zambia, Suriname, Namibia, Botswana, South Africa

Research Question 3 Is a higher tax revenue as a % of GDP associated with less income inequality?

Answer: No

Research Question 4 - Is Higher Income Per Person - GDP Per Capita associated with less income inequality?

Answer: No, but weak negative correlation.

Research Question 5 - Is Higher Investment as % GDP associated with less income inequality?

Answer: No

Research Question 6 - Is Higher EIU Democracy Index associated with less income inequality?

Answer: No, but weak negative correlation.

The above results suggest that there are other drivers for the overall reduction in income inequality. Further analysis of additional factors should be undertaken.